

Foundation Models: A Survey of the Field

Recent work on language models has shown that large-scale pretraining improves performance across many natural language processing tasks. In this paper we study the architecture, the training data, and the evaluation of foundation models on standard benchmarks.

Results and Discussion

Our analysis confirms that model size and dataset quality are the most important factors for accuracy.

The attention mechanism plays a central role in representing context, while the gradient of the loss guides the optimization of parameters during training.